

Database Systems

Introduction to Database Systems

Prof. Dr. Agnès Voisard Muhammed-Ugur Karagülle

Freie Universität Berlin,
Institute of Computer Science, Databases and Information Systems Group

Fraunhofer FOKUS

2025



- ▶ R. Elmasri and S.B. Navathe, *Fundamentals of Database Systems, 7th edition*, Benjamin/Cummings, ISBN 9781292097626, 2016
- ▶ J.D. Ullman, J. Widom, and H. Garcia-Molina, *Database Systems: Pearson New International Edition*, Pearson Education Limited, ISBN 9781292024479, 2013
- ▶ A. Silberschatz, H.F. Korth, and S. Sudarshan, *Database System Concepts Fifth Edition*, McGraw-Hill, ISBN 0072958863, 2005
- ▶ A. Kemper and A. Eickler, *Datenbanksysteme: eine Einführung*, de Gruyter Oldenbourg, ISBN 9783110443752, 2015



1 Motivation

2 Fundamentals

3 Requirements

4 Data Models & DB Schemas

5 Modern Trends

6 Challenges

7 Summary

8 Questions

9 Appendix - History



Why do we need Database Systems?

Motivation Fundamentals Requirements Data Models & DB Schemas Modern Trends Challenges Summary Questions

Imagine real-world scenarios:

- ▶ University Administration:
 - ▶ "Find all available lecture halls Friday at 10 am"
 - ▶ "List all students enrolled in the Database course"
 - ▶ "Calculate the average number of courses for the students"

Why do we need specialized database systems?

Efficient and reliable management of complex data operations



Traditional file-based systems quickly run into problems:

- ▶ Data redundancy and inconsistency
- ▶ Difficulty enforcing data integrity constraints
- ▶ Lack of concurrent access handling
- ▶ Security risks and lack of standardized access control
- ▶ Difficulty querying complex datasets efficiently

Result: Chaos, inefficiency, errors, and high maintenance costs

Databases solve these issues by providing:

- ▶ **Structured Data Representation:** Clear schema definition and metadata
- ▶ **Efficient Data Access:** Fast querying and indexing
- ▶ **Integrity and Consistency:** Reliable transactions and fault tolerance
- ▶ **Secure Access:** Controlled permissions and data protection
- ▶ **Concurrent Usage:** Safe multi-user interaction without data corruption

Database (DB) Collection of related data organized according to a schema

Database Management System (DBMS) Collection of software programs for defining, constructing, and manipulating a database

Database System (DBS) Combination of DB and DBMS software

Why separate database from software (DBMS)?

Clear separation allows flexibility, modularity, and efficient data management

A DBMS supports three core functions:

- ▶ **Defining** databases (schema specification)
 - ▶ What data will be stored?
 - ▶ What types, constraints, and structures for the data?
- ▶ **Constructing databases:** entering values, physically storing data on disk
- ▶ **Manipulating databases:**
 - ▶ Querying data to retrieve specific information
 - ▶ Modifying data (inserting, updating, deleting information)

Database systems involve different types of users:

- ▶ **End users:** Query data via applications or interfaces
- ▶ **Database Administrators (DBA):** Manage, secure, and optimize databases
- ▶ **Database designers:** Design logical structure (schema)
- ▶ **System analysts:** Identify user requirements and system specifications
- ▶ **Application programmers:** Develop software using database APIs

Practical Example

University database: Students (end users), Admin staff (DB administrators), Developers (DB designers and programmers)

Two central concepts:

► Database Schema:

- Structural description of the database
- Defines data types, structures, relationships (static)
- Metadata stored separately (self-describing)

► Database:

- Dynamic collection of data
- Changes frequently (inserts, updates, deletes)
- Represents real-world entities

- ▶ Database contains data **and metadata**:
 - ▶ **Metadata**: Description of data structures, schemas, constraints
 - ▶ Stored in a **System Catalog** managed by DBMS
- ▶ **Self-describing databases**:
 - ▶ DBS stores information about structure within itself
 - ▶ Allows software and users to interact without knowing physical details

Key Advantage: Modularity

DBMS software is generic, not tied to specific applications or data structures

What Do Database Systems Need to Provide?

Motivation Fundamentals Requirements Data Models & DB Schemas Modern Trends Challenges Summary Questions

Databases must address several key challenges:

- ▶ **Scalability:** Handle large data volumes efficiently
- ▶ **Security & Compliance:** Prohibit non-authorized access to data
- ▶ **Performance & Efficiency:** Optimize queries and processing speed
- ▶ **Reliability & Fault Tolerance:** Ensure continuous availability
- ▶ **Concurrency & Transactions:** Support multi-user access without conflicts

Why does this matter?

Every database system — SQL, NoSQL, cloud — must address these fundamental challenges



Why is database security essential?

- ▶ Prevent **unauthorized access** with authentication & encryption
- ▶ Ensure **regulatory compliance** (GDPR, HIPAA, financial standards)
- ▶ Maintain **audit logs** to track access & modifications

Without proper security...

Data breaches lead to legal and financial damage

What happens when databases grow?

- ▶ Small databases: Run on a single server
- ▶ Large-scale systems: May require **distributed architectures**
- ▶ Two common approaches:
 - ▶ **Vertical Scaling (Scale Up)**: Add more CPU/RAM to a single machine
 - ▶ **Horizontal Scaling (Scale Out)**: Distribute data across multiple servers

How do databases remain fast as they grow?

- ▶ **Indices** speed up queries
- ▶ **Query optimization** reduces processing time
- ▶ **Caching mechanisms** store frequent results

Databases must handle failures safely:

- ▶ **ACID Transactions (Atomicity, Consistency, Isolation, Durability)**
- ▶ **Backup & Recovery Strategies**
- ▶ **Replication & Redundancy** to prevent data loss

Definition

A **Data Model** defines:

- ▶ **Data Structures:** How data is organized (tables, graphs, trees)
- ▶ **Operations:** Allowed manipulations (queries, insertions, updates)
- ▶ **Constraints:** Rules ensuring integrity and correctness

Data Model is used for:

- ▶ defining the schema (Data Definition Language: DDL)
- ▶ accessing and updating the DB (Data Manipulation Language: DML)

How does a data model translate into a schema?

Database Schema (static) Defines the structure of data:

- ▶ Table definitions, data types, relationships, constraints
- ▶ Remains (relatively) stable over time

Database Instance (dynamic) Actual data:

- ▶ Frequently changing, reflecting real-world operations
- ▶ Current snapshot of database state

Database systems introduce three distinct abstraction levels:

▶ **External (Logical) Level:**

- ▶ Views tailored to specific user groups

▶ **Conceptual Level:**

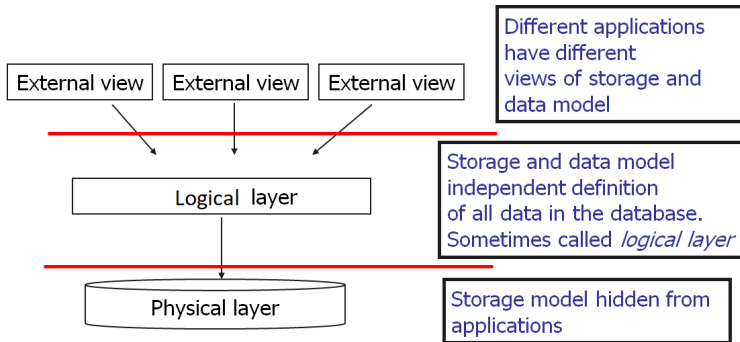
- ▶ Complete logical description (ER-model, tables)
- ▶ Independent from physical implementation

▶ **Internal (Physical) Level:**

- ▶ Actual storage structures (index, B-Trees, hashing)

Three-Schema Architecture (ANSI/SPARC) (cont'd)

Motivation Fundamentals Requirements Data Models & DB Schemas Modern Trends Challenges Summary Questions



Why three levels of abstraction?

To achieve logical and physical data independence

Database systems provide levels of abstraction to:

- ▶ Shield application programs from changes in data storage
- ▶ Allow schema modifications without affecting applications

Definition: Data Independence

The capability to change the database schema at one level without affecting the schema at a higher level

Two levels of data independence are crucial:

Physical Data Independence: Changes in the **physical schema** (e.g., indexing methods, storage devices) do not affect the logical structure or applications

Logical Data Independence: Changes to the **logical schema** (e.g., table structure) have minimal or no impact on existing applications

Goal: Reduce maintenance overhead and allow flexibility in DB evolution

Introduced by Edgar F. Codd (1970)

- ▶ Data structured as tables (relations)
- ▶ Separation of schema (metadata) and data
- ▶ Uses simple, powerful query languages (SQL)

Example: University database table

Course	Lecturer	Room
Database Systems	Prof. Voisard	T9-Gr.HS
Linear Algebra	Dr. Willert	T9-SR005
Datastructures and Algorithms	Prof. Mulzer	T9-Gr.HS

The dominating Relational DBMS

- ▶ Oracle Database
- ▶ Postgres (open-source!)
- ▶ MySQL (open-source!)
- ▶ Microsoft SQL Server
- ▶ IBM Db2
- ▶ personal, low cost desktop DBS: MSAccess

This list is not complete as the landscape is quite dynamic

**Relational data model is still the industry standard..
.. but other data models exist, which address specific needs**

Why NoSQL?

- ▶ Traditional relational databases struggle with modern scalability needs
- ▶ NoSQL databases emerged to handle large-scale, high-velocity, and semi-structured data
- ▶ Designed for flexible, distributed, and high-performance applications

Key Features of NoSQL

- ▶ Schema-less data storage
- ▶ Horizontal scalability
- ▶ More features dependent on model type

Popular NoSQL categories:

- ▶ **Key-Value Stores:** Redis, Amazon DynamoDB
- ▶ **Document Databases:** MongoDB, Couchbase
- ▶ **Column-family Stores:** Apache Cassandra, Apache HBase
- ▶ **Graph Databases:** Neo4j, Amazon Neptune

Managed database services in the cloud:

- ▶ Provides fully managed database solutions without infrastructure maintenance
- ▶ **Examples:** AWS RDS, Azure SQL Database, Google Cloud Firestore
- ▶ **Benefits:** Scalability, cost-effectiveness, automatic backups, and high availability

Realistic Example

A startup migrated its SQL database to a Cloud Database. This transition eliminated infrastructure costs and overhead while gaining automated scaling and high availability.

Processing real-time data streams:

- ▶ Used for handling **continuous, high-speed data flows**.
- ▶ **Examples:** Apache Kafka (distributed event streaming), Apache Flink (real-time analytics).
- ▶ **Benefits:** Enables **low-latency processing, event-driven architectures, and real-time decision-making**.

Realistic Example

LinkedIn uses Apache Kafka to aggregate log information from different services and systems to conduct real-time analysis of service health and performance

Decentralized data processing on the edge:

- ▶ Used for **storing and processing data closer to the source** (IoT, mobile, autonomous systems)
- ▶ **Examples:** SQLite (lightweight embedded DB), FaunaDB (distributed serverless DB)
- ▶ **Benefits:** Reduces latency, works offline, and minimizes network dependence

Realistic Example

An autonomous vehicle fleet uses an **Edge Database** to process sensor data locally. This enables **real-time navigation decisions** without relying on cloud connectivity, ensuring **faster reactions to road conditions**.

Operational requirement

- ▶ The DBS should never do anything that destroys the **consistency of database** and modeled reality (called **integrity**)
- ▶ **Main technical issue:** Execution of operations must **guarantee correctness properties**

Example

Transfer \$ 100 from one account a1 to another one a2.

- 1 Reading the value x of account a1
- 2 Decrease the value x by \$ 100
- 3 Write the new value of x to the account a1
- 4 Read the value y of account a2
- 5 Increase the value of y by \$ 100
- 6 Write the new value of y to the account a2

Operational requirement

- ▶ **No interference of operations** of different users

Example

Reservation system: Two independent users want to reserve the same seat on a plane

Fail-safe operation

- ▶ **System failure** should **not corrupt database** state

Example

System crash when writing new account balance on disk.
DB must not be corrupted

Efficiency

- ▶ Hundreds of clients active on the same DB
- ▶ Hundreds or thousands operations / sec
- ▶ Response time requirement in interactive environment: < 3 sec

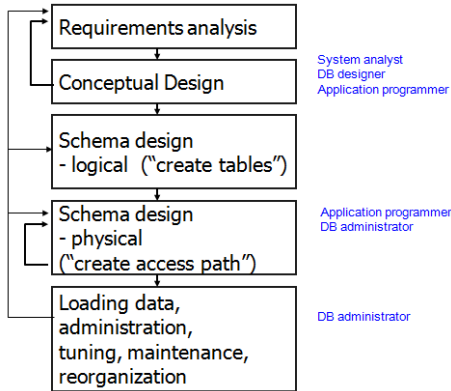
Data security

- ▶ Access by unauthorized users might be a disaster

Synchronisation of independent DB-users

- ▶ How to avoid conflicting read / write access?
⇒ concurrent programming
- ▶ But DB have many resources:
Each record is a resource – there may be millions¹ of them!
⇒ Synchronization of thousands of concurrent operations?

¹Wal-Mart: 200 mio transaction / week = 300 TA/sec (24/7)



Compare

► Life cycle of hardware is about 3 years

► Life cycle of software is about 5 years

► Life cycle of data is about 30 years

- ▶ Database \neq Database System
- ▶ Database
 - ▶ **Data**
 - ▶ **Metadata** (Schema)
- ▶ Data model
- ▶ Relational Data Model (RDM) / SQL
- ▶ NoSQL Data Models
- ▶ Technical Requirements
 - ▶ Concurrency
 - ▶ Fault-tolerance
 - ▶ Integrity
 - ▶ Efficiency
- ▶ New Trends Challenges
- ▶ Lifecycle

Questions?

Motivation Fundamentals Requirements Data Models & DB Schemas Modern Trends Challenges Summary Questions



What will come next?

Motivation Fundamentals Requirements Data Models & DB Schemas Modern Trends Challenges Summary Questions

- 1 Welcome to Database Systems
- 2 Introduction to Database Systems
- 3 Entity Relationship Design Diagram (ERM)
- 4 Relational Model
- 5 Relational Algebra
- 6 Structured Query Language (SQL)
- 7 Relational Database Design - Functional Dependencies
- 8 Relational Database Design - Normalization
- 9 Online Analytical Processing + Embedded SQL
- 10 ~~Data Mining~~
- 11 Physical Representation - Storage and File Structure
- 12 Physical Representation - Indexing and Hashing
- 13 Transactions
- 14 Concurrency Control Techniques
- 15 Recovery Techniques
- 16 Query Processing and Optimization



- ▶ Business Data Processing as the driving force for DBS development
- ▶ about 1965 file system approach to data management leads to chaos
- ▶ **What are the right abstractions? \Rightarrow Data model**
- ▶ 1970: Tables!
- ▶ 1973: Research prototypes for Relational DBS, Transactions
- ▶ 1980: RDBMS everywhere, Distributed DBS

- ▶ 1990: Object orientation \Rightarrow OO data model and OODBMS \Rightarrow Object-Relational systems
- ▶ 1995: Wide scale distribution, **WEB**
- ▶ 1997: Semistructured data, Image DB, ..., XML / DB
- ▶ 2000++ Mobility and DBMS
- ▶ 2005++ Unstructured Data – e.g., text. Querying text?
- ▶ Automated **Object-relational mapping**:
Only objects in the program, relations are not the main focus