

# Introduction to DBMS Transaction Processing

**Prof. Dr. Agnès Voisard**

Institute of Computer Science Databases and  
Information Systems Group  
and Fraunhofer FOKUS

**Summer term 2025**  
**v2**

# Contents

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability

# Read/Write operations

- Level of data items and disk blocks. Data item: field of a DB record, record, whole block
- DB operations
  - **Read-item (X) or read (X).** Execution:
    - ▶ Find the address of the disk block that contains X
    - ▶ Copy the disk block into a buffer in main memory
    - ▶ Copy item X from the buffer to the program variable
  - **Write-item (X) or write (X).** Execution:
    - ▶ Find the address of the disk block that contains X
    - ▶ Copy the disk block into a buffer
    - ▶ Copy item X from the program variable X into its correct place in the buffer
    - ▶ Store the updated block from the buffer back to disk (update of DB on disk)

# Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items
- E.g., transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- Two main issues to deal with:
  - Failures of various kinds, such as hardware failures and system crashes
  - Concurrent execution of multiple transactions

# Example of Fund Transfer

- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- **Atomicity requirement**
  - If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
    - ▶ Failure could be due to software or hardware
  - The system should ensure that updates of a partially executed transaction are not reflected in the database
- **Durability requirement** Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures

# Example of Fund Transfer (cont'd)

- Transaction to transfer \$50 from account A to account B:
  1. **read**(A)
  2.  $A := A - 50$
  3. **write**(A)
  4. **read**(B)
  5.  $B := B + 50$
  6. **write**(B)
- **Consistency requirement** in above example:
  - The sum of A and B is unchanged by the execution of the transaction
- In general, consistency requirements include
  - ▶ Explicitly specified integrity constraints such as primary keys and foreign keys
  - ▶ Implicit integrity constraints
    - e.g. , sum of balances of all accounts, minus sum of loan amounts must equal the value of cash-in-hand
  - **A transaction must see a consistent database.**
  - During transaction execution the database may be temporarily inconsistent.
  - When the transaction completes successfully the database must be consistent
    - ▶ Erroneous transaction logic can lead to inconsistency

# Example of Fund Transfer (cont'd)

- **Isolation requirement** If between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be)

**T1**

1. **read**(A)
2.  $A := A - 50$
3. **write**(A)
4. **read**(B)
5.  $B := B + 50$
6. **write**(B)

**T2**

read(A), read(B), print(A+B)     **Dirty read**

- Isolation can be ensured trivially by running transactions **serially**
  - I.e., one after the other
  - However, executing multiple transactions concurrently has significant benefits

# ACID Properties

A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:

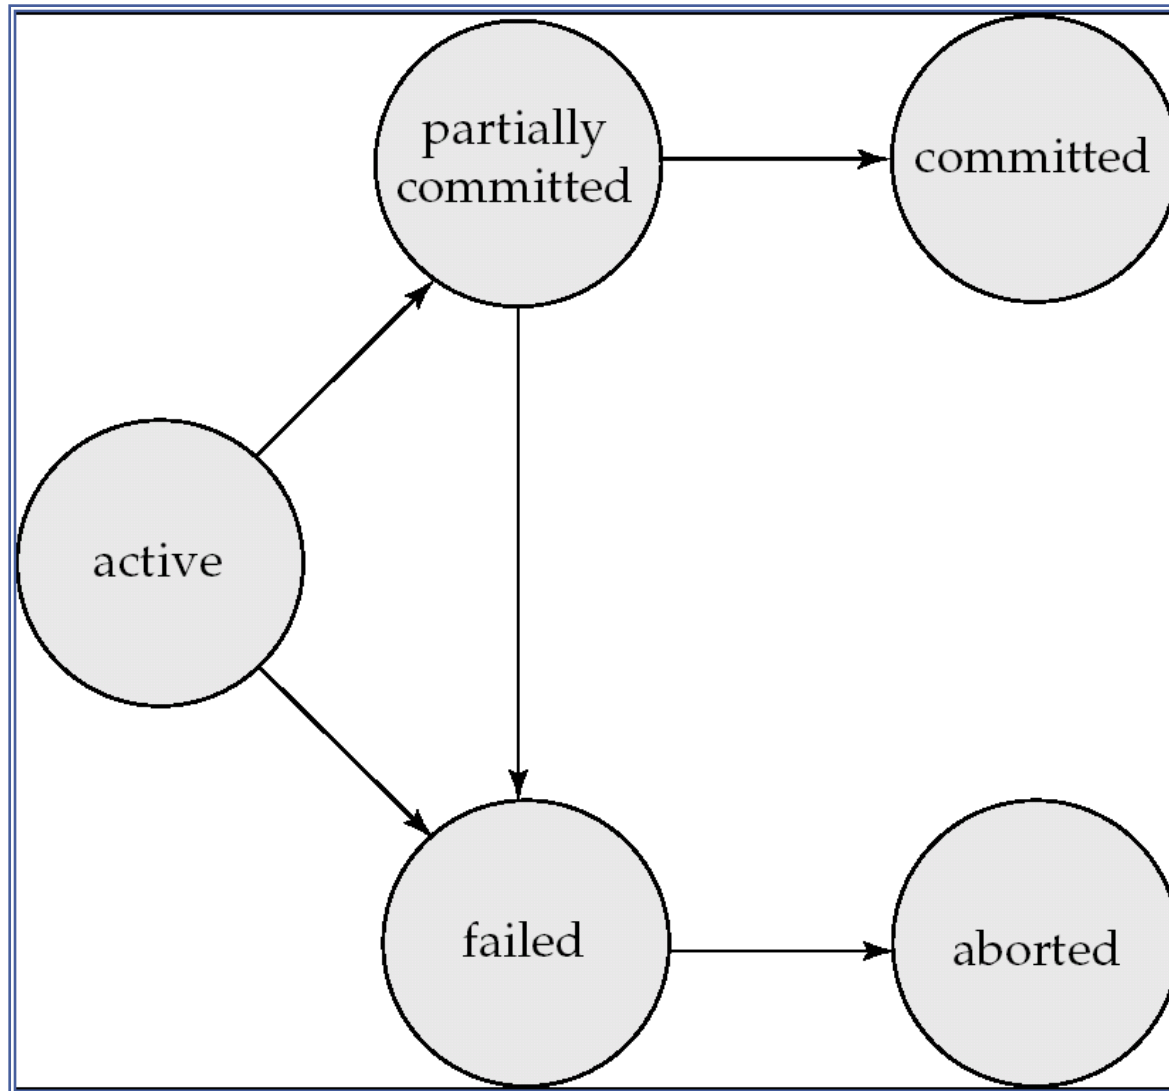
- **Atomicity** Either all operations of the transaction are properly reflected in the database or none are
- **Consistency** Execution of a transaction in isolation preserves the consistency of the database
- **Isolation** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions
  - That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished
- **Durability** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures



# Transaction State

- **Active** – The initial state; the transaction stays in this state while it is executing
- **Committed** – After successful completion
- **Partially committed** – After the final statement has been executed
- **Failed** – After the discovery that normal execution can no longer proceed
- **Aborted** – After the transaction has been rolled back and the database restored to its state prior to the start of the transaction. Two options after it has been aborted:
  - Restart the transaction
  - Error kill the transaction

# Transaction State (cont'd)



# Introduction to recovery

- Transaction can fail in the middle of execution
  
- Types of failure:
  - System crash
  - Transaction or system error (e.g., division by 0)
  - Errors detected during the transaction (e.g., data not found)
  - Concurrency control enforcement (transaction has to be aborted bc of deadlocks, see later)
  - Disk failure (blocks lose their data. May happen during a read/write operation)
  - Physical problems and catastrophes: power failure, sabotage, etc.

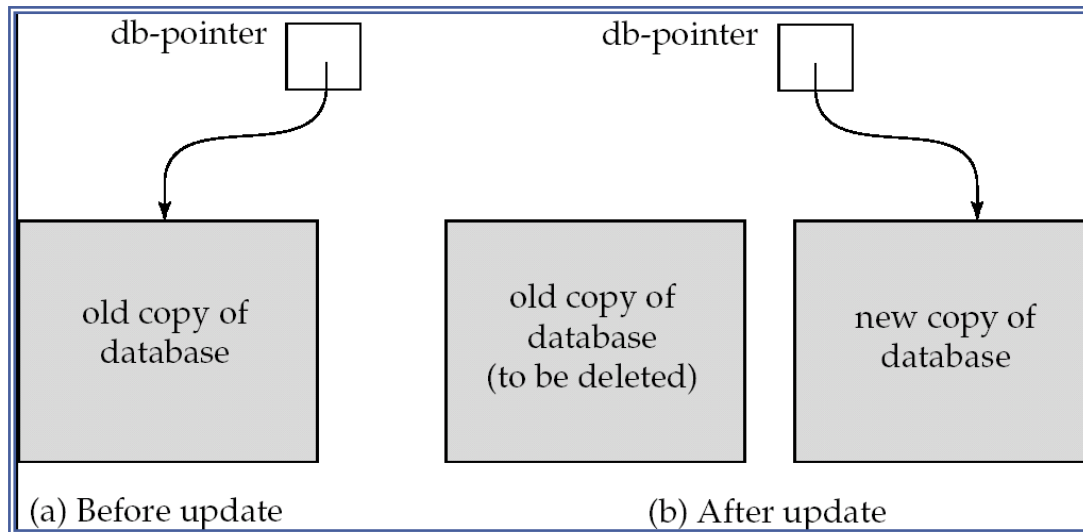
# Introduction to recovery (cont'd)

Recovery manager keeps track of when the transaction starts, terminates, commits, or abort:

- BEGIN TRANSACTION
  - READ or WRITE DB item
  - COMMIT-TRANSACTION
    - Signals a successful end of the transaction
  - ROLL-BACK (or ABORT)
    - Transaction has ended unsuccessfully
    - Any change done by the transaction must be undone
  - END-TRANSACTION: READ and WRITE have ended
    - End of the transaction BUT it must be checked:
      - ▶ Whether changes can be permanently applied to the DB (COMMIT) or
      - ▶ Whether the transaction has to be aborted
- + UNDO/REDO, see later
- This is written in a journal (log)

# Implementation of Atomicity and Durability

- The **recovery-management** component of a database system implements the support for atomicity and durability
- E.g., the *shadow-database* scheme:
  - All updates are made on a *shadow copy* of the database
    - ▶ **db\_pointer** is made to point to the updated shadow copy after
      - The transaction reaches **partial commit** and
      - All updated pages have been flushed to disk.



# Implementation of Atomicity and Durability (cont'd)

- db\_pointer always points to the current consistent copy of the database
  - In case transaction fails, old consistent copy pointed to by **db\_pointer** can be used, and the shadow copy can be deleted
- The shadow-database scheme:
  - Assumes that only one transaction is active at a time
  - Assumes disks do not fail
  - Useful for text editors, but
    - ▶ Extremely inefficient for large databases
    - ▶ Variant called shadow paging reduces copying of data, but is still not practical for large databases
  - Does not handle concurrent transactions

# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system  
Advantages are:
  - **Increased processor and disk utilization**, leading to better transaction *throughput*
    - ▶ E.g., one transaction can be using the CPU while another is reading from or writing to the disk
  - **Reduced average response time** for transactions:  
short transactions need not wait behind long ones
- **Concurrency control schemes**: Mechanisms to achieve isolation
  - I.e., to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database

# Schedules

- **Schedule:** A sequence of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - A schedule for a set of transactions must consist of all instructions of those transactions
  - A schedule must preserve the order in which the instructions appear in each individual transaction
- A transaction that successfully completes its execution will have a **commit** instruction as the last statement
  - By default, a transaction is assumed to execute commit instruction as its last step
- A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement



# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$  and  $T_2$  transfer 10% of the balance from  $A$  to  $B$
- A **serial** schedule in which  $T_1$  is followed by  $T_2$ :

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write ( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ ) $B := B + temp$ write( $B$ )

# Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ ) $B := B + temp$ write( $B$ )

# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1

$T_1$	$T_2$
read(A) $A := A - 50$ write(A)	read(A) $temp := A * 0.1$ $A := A - temp$ write(A)
read(B) $B := B + 50$ write(B)	read(B) $B := B + temp$ write(B)

In Schedules 1, 2, and 3, the sum  $A + B$  is preserved

# Schedule 4

- The following concurrent schedule does not preserve the value of  $(A + B)$

$T_1$	$T_2$
read( $A$ ) $A := A - 50$	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ )
write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	$B := B + temp$ write( $B$ )

# Serializability

- **Basic Assumption:** Each transaction preserves database consistency
- Thus serial execution of a set of transactions preserves database consistency
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
  1. **Conflict serializability**
  2. **View serializability**
- *Simplified view of transactions*
  - We ignore operations other than **read** and **write** instructions
  - We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes
  - Our simplified schedules consist of only **read** and **write** instructions

# Conflicting Instructions

- Instructions  $l_i$  and  $l_j$  of transactions  $T_x$  and  $T_y$  respectively, **conflict** if and only if there exists some **item  $Q$  accessed by both  $l_x$  and  $l_y$ , and at least one of these instructions wrote  $Q$ .**
  1.  $l_x = \text{read}(Q)$ ,  $l_y = \text{read}(Q)$ .  $l_x$  and  $l_y$  don't conflict
  2.  $l_x = \text{read}(Q)$ ,  $l_y = \text{write}(Q)$ . They conflict
  3.  $l_x = \text{write}(Q)$ ,  $l_y = \text{read}(Q)$ . They conflict
  4.  $l_x = \text{write}(Q)$ ,  $l_y = \text{write}(Q)$ . They conflict
- Intuitively, a conflict between  $l_x$  and  $l_y$  forces a (logical) temporal order between them
  - If  $l_x$  and  $l_y$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule

# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule

# Conflict Serializability (cont'd)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions
- Therefore Schedule 3 is conflict serializable

$T_1$	$T_2$
read( $A$ ) write( $A$ )	read( $A$ ) write( $A$ )
read( $B$ ) write( $B$ )	
	read( $B$ ) write( $B$ )

Schedule 3

$T_1$	$T_2$
read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )	read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )

Schedule 6



# Conflict Serializability (cont'd)


- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read( $Q$ )	write( $Q$ )
write( $Q$ )	

- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$

# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph**: a direct graph where the vertices are the transactions (names)
- We draw an arc from  $T_x$  to  $T_y$  if the two transaction conflict, and  $T_x$  accessed the data item on which the conflict arose earlier
- We may label the arc by the item that was accessed.
- **Example 1**




# Precedence graph - constructionAlgorithm

- Given: schedule S.
- 1. For each  $T_x$  in S create a node
- 2. For each case in S where  $T_y$  does a read (X) after a write (X) by  $T_x$ 
  - $T_x$ : write (X) ;  $T_y$ : read (X) ;
  - Create ( $T_x \rightarrow T_y$ )
- 3. For cases where  $T_y$  does a write (X) after  $T_x$  does a read (X)
  - $T_x$ : read (X) ;  $T_y$ : write (X) ;
  - Create ( $T_x \rightarrow T_y$ )
- 4. For cases where  $T_y$  does a write (X) after  $T_x$  does a write (X)
  - $T_x$ : write (X) ;  $T_y$ : write (X) ;
  - Create ( $T_x \rightarrow T_y$ )
- 5. S is serializable if the precedence graph has no cycle

## Example Schedule (Schedule A) + Precedence Graph

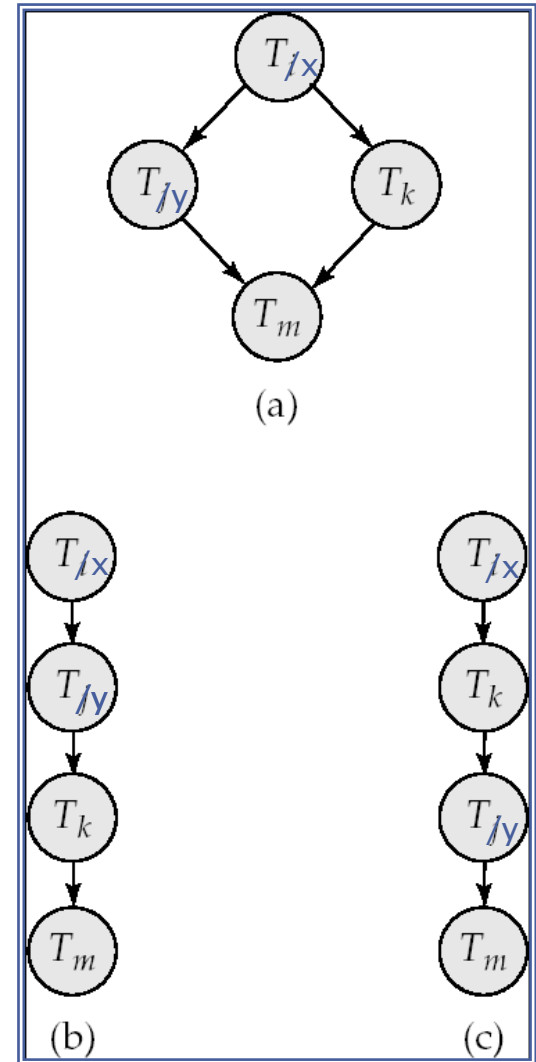
$T_1$	$T_2$	$T_3$	$T_4$	$T_5$
read(Y) read(Z)	read(X)			read(V) read(W) read(W)
	read(Y) write(Y)	write(Z)		
read(U)			read(Y) write(Y) read(Z) write(Z)	
read(U) write(U)				



$T_5$

# Test for Conflict Serializability

- A schedule is **conflict** serializable if and only if its precedence graph is acyclic
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph  
*Better algorithms take order  $n + e$  where  $e$  is the number of edges*
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph
  - This is a linear order consistent with the partial order of the graph
  - For example, a serializability order for Schedule (a) would be  
 $T_5 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2 \rightarrow T_4$




# Simple notation

□ S, transactions T1, T2, T3, items u, x, y

□ S: r1(y), r3(u), r2(y), w1(y), w2(x), w1(x), w2(z), w3(x)

□ Precedence graph



Serializable!

# Concurrency Control

- A database must provide a mechanism that will ensure that all possible schedules are
  - Either **conflict** or **view serializable**, and
  - Are recoverable and preferably cascadeless (see later)
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- Testing a schedule for serializability *after* it has executed is a little too late!
- **Goal:** To develop concurrency control protocols that will assure serializability

# Concurrency Control vs. Serializability Tests

- Concurrency-control protocols allow concurrent schedules but ensure that the schedules are conflict/view serializable, and are recoverable and cascadeless
- Concurrency control protocols generally **do not examine the precedence graph** as it is being created
  - Instead, a protocol imposes a discipline that **avoids non-serializable schedules**
- Different concurrency control protocols provide different tradeoffs between the amount of concurrency they allow and the amount of overhead that they incur
- Tests for serializability help us understand why a concurrency control protocol is correct



# Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - E.g., a **read-only transaction** that wants to get an approximate total balance of all accounts
  - E.g., database statistics computed for query optimization can be approximate
  - Such transactions need not be serializable with respect to other transactions
- Tradeoff **accuracy** for **performance**

# Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction
- In SQL, a transaction begins implicitly
- A transaction in SQL ends by:
  - **Commit work** commits current transaction and begins a new one
  - **Rollback work** causes current transaction to abort
- In almost all database systems, by default, every SQL statement also **commits implicitly if it executes successfully**
  - Implicit commit can be turned off by a database directive
    - ▶ E.g., in JDBC, `connection.setAutoCommit(false);`

# Summary

- Transaction concept
- ACID properties
- Introduction to recovery
- Serial schedules
- Conflict (precedence) graph



# Levels of Consistency in SQL-92

- **Serializable** – default
  - **Repeatable read** – only committed records to be read, **repeated reads of same record must return same value**. However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others
  - **Read committed** – only committed records can be read, but **successive reads of record may return different** (but committed) values
  - **Read uncommitted** – even uncommitted records may be read
- 
- Lower degrees of consistency useful for gathering approximate information about the database
  - Warning: some database systems do not ensure serializable schedules by default
    - E.g., Oracle and PostgreSQL by default support a level of consistency called snapshot isolation (not part of the SQL standard)